

## ▼ Diabetes Prediction — Using a Local CSV

This notebook uses a local CSV file (`pima_diabetes_dataset.csv`) (placed in the same folder). It performs a full ML workflow with explanations:

- Load local CSV
- Inspect data and handle missing values
- Preprocess (replace invalid zeros, impute using median, scale features)
- Train 5 classifiers (Logistic Regression, Decision Tree, Random Forest, KNN, SVC)
- Evaluate models (accuracy, ROC-AUC, classification report, confusion matrix)
- Plot model comparison and give short assignment notes

Run the notebook cells top-to-bottom. The notebook is written for beginners with comments and brief explanations.

### ▼ 1) Imports and helper functions

```
# Imports
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, roc_auc_score
import matplotlib.pyplot as plt

print('Libraries imported')
```

Libraries imported

### ▼ 2) Load the local CSV dataset

```
# Load local CSV (ensure the file is in the same folder as this notebook)
csv_path = 'pima_diabetes_dataset.csv' # change path if needed
df = pd.read_csv(csv_path)
print('Loaded shape:', df.shape)
df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
0	6	148	72	35	0	33.6		0.627	50	1
1	1	85	66	29	0	26.6		0.351	31	0
2	8	183	64	0	0	23.3		0.672	32	1
3	1	89	66	23	94	28.1		0.167	21	0
4	0	137	40	35	168	43.1		2.288	33	1

### ▼ 3) Basic inspection

```
print(df.info())
print('\nSummary statistics:\n', df.describe().T)
print('\nTarget value counts:\n', df['Outcome'].value_counts())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype  
 ---  -- 
 0   Pregnancies      10 non-null    int64  
 1   Glucose          10 non-null    int64  
 2   BloodPressure    10 non-null    int64  
 3   SkinThickness    10 non-null    int64  
 4   Insulin          10 non-null    int64  
 5   BMI              10 non-null    float64 
 6   DiabetesPedigreeFunction 10 non-null    float64 
 7   Age              10 non-null    int64  
 8   Outcome          10 non-null    int64 
```

```

0 Pregnancies          10 non-null    int64
1 Glucose              10 non-null    int64
2 BloodPressure        10 non-null    int64
3 SkinThickness        10 non-null    int64
4 Insulin              10 non-null    int64
5 BMI                  10 non-null    float64
6 DiabetesPedigreeFunction 10 non-null float64
7 Age                  10 non-null    int64
8 Outcome              10 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 852.0 bytes
None

Summary statistics:
   count      mean       std      min     25%  \
Pregnancies      10.0    4.4000  3.502380  0.000  1.2500
Glucose          10.0  127.3000 40.075068 78.000 95.5000
BloodPressure    10.0    59.8000 25.654976  0.000 53.5000
SkinThickness    10.0   19.9000 17.978073  0.000  0.0000
Insulin          10.0   89.3000 169.937276  0.000  0.0000
BMI               10.0   27.7100 11.259016  0.000 25.8500
DiabetesPedigreeFunction 10.0    0.5078  0.654114  0.134  0.1755
Age              10.0   35.9000 11.873874 21.000 29.2500
Outcome          10.0    0.6000  0.516398  0.000  0.0000

   50%      75%      max
Pregnancies      4.00    7.500 10.000
Glucose          120.50 145.250 197.000
BloodPressure    66.00   71.500 96.000
SkinThickness    26.00   34.250 45.000
Insulin          0.00   92.500 543.000
BMI              29.30   32.950 43.100
DiabetesPedigreeFunction 0.24   0.558  2.288
Age              31.50   45.750 54.000
Outcome          1.00    1.000 1.000

Target value counts:
  Outcome
  1    6
  0    4
Name: count, dtype: int64

```

#### ▼ 4) Robust target handling (ensure y is numeric 0/1)

```

# Robust target mapping (handles string labels if present)
y = df['Outcome']
import numpy as np
if not np.issubdtype(pd.Series(y).dtype, np.number):
    print('Non-numeric target detected. Examples:', pd.Series(y).unique()[:20])
    mapping = {
        'tested_negative': 0, 'tested_positive': 1,
        'negative': 0, 'positive': 1,
        'No': 0, 'Yes': 1,
        'no': 0, 'yes': 1,
        'NEGATIVE': 0, 'POSITIVE': 1,
        '0': 0, '1': 1
    }
    y_mapped = pd.Series(y).map(mapping)
    if y_mapped.isna().any():
        coerced = pd.to_numeric(pd.Series(y), errors='coerce')
        if coerced.isna().any():
            missing = pd.Series(y)[y_mapped.isna()].unique()
            raise ValueError(f'Cannot automatically map these target labels: {missing}. Add them to mapping.')
        else:
            y = coerced.astype(int)
    else:
        y = y_mapped.astype(int).values
else:
    y = pd.Series(y).astype(int).values

print('Final target unique values:', np.unique(y))

Final target unique values: [0 1]

```

#### ▼ 5) Preprocessing: replace invalid zeros, impute, and scale

```
# Features
X = df.drop(columns=['Outcome']).copy()

# Columns where 0 is medically invalid and likely represents missing values
cols_with_zero_invalid = ['Glucose', 'BloodPressure', 'SkinThickness', 'Insulin', 'BMI']

# Replace zeros with NaN in those columns
X[cols_with_zero_invalid] = X[cols_with_zero_invalid].replace(0, np.nan)
print('Missing counts after replacing zeros:')
print(X[cols_with_zero_invalid].isna().sum())

# Impute missing values with median
imputer = SimpleImputer(strategy='median')
X_imputed = pd.DataFrame(imputer.fit_transform(X), columns=X.columns)

# Scale features
scaler = StandardScaler()
X_scaled = pd.DataFrame(scaler.fit_transform(X_imputed), columns=X_imputed.columns)

print('\nAfter imputation and scaling - feature summary:')
print(X_scaled.describe().T)

Missing counts after replacing zeros:
Glucose      0
BloodPressure 1
SkinThickness 4
Insulin       6
BMI          1
dtype: int64

After imputation and scaling - feature summary:
      count      mean      std      min      25%  \
Pregnancies  10.0 -8.881784e-17  1.054093 -1.324244 -0.948039
Glucose      10.0  8.326673e-17  1.054093 -1.296735 -0.836434
BloodPressure 10.0 -3.552714e-16  1.054093 -1.890349 -0.136048
SkinThickness 10.0  5.780099e-16  1.054093 -1.985548 -0.178314
Insulin      10.0 -5.551115e-17  1.054093 -0.630258 -0.291071
BMI          10.0 -1.859624e-16  1.054093 -1.390562 -0.705533
DiabetesPedigreeFunction 10.0 -2.498002e-16  1.054093 -0.602372 -0.535495
Age          10.0  1.110223e-16  1.054093 -1.322734 -0.590348

      50%      75%      max
Pregnancies -0.120386  0.932990  1.685402
Glucose     -0.178860  0.472138  1.833316
BloodPressure -0.028642  0.365181  2.119483
SkinThickness  0.038554  0.255422  2.255428
Insulin      -0.291071 -0.291071  2.958823
BMI          -0.048465  0.408221  2.300206
DiabetesPedigreeFunction -0.431555  0.080896  2.868759
Age          -0.390606  0.874425  1.606811
```

## 6) Train/Test split

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.30, random_state=42, stratify=y)
print('Train shape:', X_train.shape)
print('Test shape:', X_test.shape)
print('Train class distribution:\n', pd.Series(y_train).value_counts())

Train shape: (7, 8)
Test shape: (3, 8)
Train class distribution:
1    4
0    3
Name: count, dtype: int64
```

## 7) Train 5 models and evaluate

```
models = {
    'LogisticRegression': LogisticRegression(max_iter=1000, random_state=42),
    'DecisionTree': DecisionTreeClassifier(random_state=42),
    'RandomForest': RandomForestClassifier(n_estimators=100, random_state=42),
    'KNN': KNeighborsClassifier(n_neighbors=5),
    'SVC': SVC(probability=True, random_state=42)
}
```

```
results = []
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    try:
        roc = roc_auc_score(y_test, model.predict_proba(X_test)[:,1])
    except Exception:
        roc = None
    results.append({'model': name, 'accuracy': acc, 'roc_auc': roc})
print(f'--- {name} ---')
print('Accuracy:', round(acc, 4))
print('Classification report:\n', classification_report(y_test, y_pred))
print('Confusion matrix:\n', confusion_matrix(y_test, y_pred))
print('\n')

results_df = pd.DataFrame(results).sort_values('accuracy', ascending=False).reset_index(drop=True)
print('Summary results:')
results_df
```




```
--- LogisticRegression ---
```

```
Accuracy: 0.6667
```

```
Classification report:
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

```
import matplotlib.pyplot as plt
plt.figure(figsize=(8,4))
plt.bar(results_df['model'], results_df['accuracy'])
plt.title('Model accuracy comparison')
plt.ylabel('Accuracy')
plt.ylim(0,1)
for i, v in enumerate(results_df['accuracy']):
    plt.text(i, v+0.01, f'{v:.3f}', ha='center')
plt.show()
```



```
MetricWarning: Precision is ill-defined
```

```
MetricWarning: Precision is ill-defined
```

```
/usr/local/lib/python3.12/dist-packages/scikit-learn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined
```

```
_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

## 9) Assignment notes and suggestions